

Contents

1	README	2
2	CDE.h	3
3	CDE.cpp	4
4	CachingFileSystem.cpp	5
5	CountChain.h	17
6	CountChain.cpp	18
7	LRUStack.h	20
8	LRUStack.cpp	21
9	Makefile	23

1 README

```
1  matanmo, eyalc
2  Matan Modai (305097560), Eyal Cohen (301300372)
3  EX: 4
4
5  FILES:
6  CachingFileSystem.cpp - main caching file
7  CDE.h - CDE CLASS
8  CDE.cpp
9  CountChain.h - Count Chain Class
10 CountChain.cpp
11 LRUStack.h - LRU Stack class
12 LRUStack.cpp
13 Makefile
14 README - this file
15
16
17
18 REMARKS:
19
20 Design and implementation is exactly as explained in chapter 2.4 in the article:
21 "Data Cache Management Using Frequency-Based Replacement"
22
23
24 ANSWERS:
25
26 1.
27 ==
28 Usually using the cache will be faster, but sometimes the OS saves some of the processs pages on
29 the disk and then when the process tries to retrieve the data, a page fault is encountered. In that
30 case, reading the data from the cache involves disk I/O and therefore the access speed is not
31 improved.
32
33 2.
34 ==
35 When using simpler algorithms, the logic for managing the cache can be implemented on hardware
36 level, therefore leaving the CPU free to perform other tasks. However, when using more sophisticated
37 algorithms, the logic is performed at the CPU level, therefore making the pages management harder.
38
39 3.
40 ==
41 When cache size is 3:
42 A,A,A,A,B,C,D,A,A,A,A,E,F,G,A,A,A,A - LFU is better than LRU because clearly A is used
43
44 A,A,A,A,A,B,C,D,E,B,C,D,E - LRU is better than LFU becace A is used at high frequency at first but
45 not at all afterwords.
46
47 A,B,C,D,A,B,C,D,A,B,C,D - MRU would be better than LRU and FRU
48
49 4.
50 ==
51 Sometimes when using resources, a resource can have high locality frequency, meaning the resource
52 will be used a lot in a brief section of the program (e.g loops), but not very much used afterwords.
53 In those cases, we wouldnt want to give this resource too much significance because the frequency
54 is only relevant for a short section - high locality frequency.
55 When using the new section, we can have a better sense if the frequency of the resource is local.
56 Using this system, a high frequency count shows that the frequency does not refer only to a small
57 section, and therefore the probability of needing the resource in the future is higher than if it
58 the frequency was local.
```

2 CDE.h

```
1  #ifndef OS_EX4_CDE_H
2  #define OS_EX4_CDE_H
3
4  #include <string>
5  #include <cstdlib>
6  #include <stdlib.h>
7
8  using namespace std;
9
10 class CDE {
11 private:
12     int _count;
13     CDE * _countPrev;
14     CDE * _countNext;
15     CDE * _prev;
16     CDE * _next;
17     bool _isOld, _isNew;
18     int _blockId;
19     string _fileName;
20     char * _blockData;
21     size_t _dataSize;
22 public:
23     CDE(int blockId, string fileName, size_t dataSize, char * blockData);
24     ~CDE();
25
26     char * getData() { return _blockData; };
27     size_t getSize() { return _dataSize; };
28
29     void setPrev(CDE * cde) { _prev = cde; };
30     void setNext(CDE * cde) { _next = cde; };
31
32     void setCountPrev(CDE * cde) { _countPrev = cde; };
33     void setCountNext(CDE * cde) { _countNext = cde; };
34
35     CDE * getPrev() { return _prev; };
36     CDE * getNext() { return _next; };
37
38     CDE * getCountPrev() { return _countPrev; };
39     CDE * getCountNext() { return _countNext; };
40
41     void increaseCount() { ++_count; };
42
43     void setIsNew(bool b) { _isNew = b; };
44     void setIsOld(bool b) { _isOld = b; };
45
46     bool getIsOld() { return _isOld; };
47     bool getIsNew() { return _isNew; };
48
49     void setFileName(string fileName) { _fileName = fileName; };
50
51     int getCount() { return _count; };
52
53     int getBlockId() { return _blockId; };
54     string getFileName() { return _fileName; };
55 };
56
57
58 #endif //OS_EX4_CDE_H
```

3 CDE.cpp

```
1  #include <cstring>
2  #include "CDE.h"
3  #include <iostream>
4
5  CDE::CDE(int blockId, string fileName, size_t dataSize,
6          char * blockData) {
7      _dataSize = dataSize;
8      _fileName = fileName;
9      _blockId = blockId;
10     _blockData = (char *) malloc(sizeof(char) * dataSize);
11     if(_blockData == NULL) {
12         cout << "malloc ERROR" << endl;
13     }
14     memcpy(_blockData, blockData, dataSize);
15     _isOld = false;
16     _isNew = true;
17     _count = 1;
18     _countNext = nullptr;
19     _countPrev = nullptr;
20     _next = nullptr;
21     _prev = nullptr;
22 };
23
24 CDE::~CDE() {
25     free(_blockData);
26 };
```

4 CachingFileSystem.cpp

```
1  /*
2   * CachingFileSystem.cpp
3   *
4   * Author: Netanel Zakay, HUJI, 67808 (Operating Systems 2015-2016).
5   */
6
7  #define FUSE_USE_VERSION 26
8
9  #include <iostream>
10 #include <fuse.h>
11 #include <unistd.h>
12 #include <sstream>
13 #include <fstream>
14 #include <limits.h>
15 #include <errno.h>
16 #include <cstdlib>
17 #include <dirent.h>
18 #include <cstring>
19 #include <map>
20 #include <cmath>
21 #include <vector>
22
23 #include "CDE.h"
24 #include "LRUStack.h"
25 #include "CountChain.h"
26
27 using namespace std;
28
29 #define CMAX 3 // the maximum count list value
30 #define USAGE_ERROR "Usage: CachingFileSystem " \
31     "rootdir mountdir numberOfBlocks fOld fNew\n"
32 #define FILE_LOCATION "/.filesystem.log"
33
34 struct fuse_operations caching_oper;
35
36 struct {
37     char* rootDir;
38 } typedef user_data;
39
40 static char* rootDir;
41 static char logPath[PATH_MAX];
42 static int blockSize;
43 static map<pair<string, int>, CDE*> cacheMap;
44 static LRUStack lru;
45 static CountChain countChain(CMAX);
46 static int numberOfBlocks;
47
48 std::ofstream logFile;
49 std::stringstream logBuffer;
50
51 /* Return the System block size*/
52 static int getBlockSize() {
53     struct stat fi;
54     stat("/tmp", &fi);
55     return (int) fi.st_blksize;
56 }
57
58 static bool isLogFile(const char *path) {
59     string logFile(".filesystem.log");
```

```

60     string str(path);
61     int pos = str.find(logFile);
62     if(pos == -1) {
63         return false;
64     }
65     return true;
66 }
67
68 /* Build absolute path from relative */
69 static void buildPath(char fpath[PATH_MAX], const char *path) {
70     strcpy(fpath, rootDir);
71     strncat(fpath, path, PATH_MAX);
72 }
73
74 /*
75  * Check that the given parameters are valid
76  */
77 bool isInputParamsValid(int argc, char* argv[]) {
78     // Usage error
79     bool usage_error = false;
80     if(argc != 6) {
81         usage_error = true;
82     } else {
83         struct stat sb;
84         // fetch params
85         char* rootdir = argv[1];
86         char* mountdir = argv[2];
87         int numberOfBlocks = atoi(argv[3]);
88         double fOld = atof(argv[4]);
89         double fNew = atof(argv[5]);
90
91         // check params are valid
92         if(fOld <= 0 || fOld >= 1 || fNew <= 0 ||
93            fNew >= 1 || fOld + fNew > 1) {
94             usage_error = true;
95         } else if(numberOfBlocks < 0) {
96             usage_error = true;
97         } else if(stat(rootdir, &sb) != 0 || !S_ISDIR(sb.st_mode)) {
98             usage_error = true;
99         } else if(stat(mountdir, &sb) != 0 || !S_ISDIR(sb.st_mode)) {
100             usage_error = true;
101         }
102     }
103     if(usage_error) {
104         cout << USAGE_ERROR << endl;
105     }
106     return !usage_error;
107 }
108
109 /*
110  * Print system errors to cerr.
111  */
112 void sysError(std::string errFunc) {
113     std::cerr << "System Error: " <<
114     errFunc << " failed." << std::endl;
115     exit(1);
116 }
117
118 /**
119  * Write commands to log.
120  */
121 void writeToLog(string msg) {
122     // logFile.open(logPath, std::ios_base::app);
123     if(logFile.fail()) {
124         sysError("open");
125     }
126
127     time_t t;

```

```

128     if((int) time(&t) < 0) {
129         sysError("time");
130     }
131
132     logFile << t << " " << msg << "\n";
133     // logFile.close();
134     if(logFile.fail()) {
135         sysError("close");
136     }
137 }
138
139
140 /**
141  * Remove block from cache
142  */
143 void removeFromCache() {
144     CDE * cde = countChain.getItemToRemove();
145     if(cde == nullptr) {
146         cde = lru.getTail();
147     }
148     lru.remove(cde);
149     countChain.remove(cde);
150     string fileName = cde->getFileName();
151     int blockId = cde->getBlockId();
152     cacheMap.erase({fileName, blockId});
153     delete cde;
154 }
155
156 /*****
157
158  /** Get file attributes.
159  *
160  * Similar to stat(). The 'st_dev' and 'st_blksize' fields are
161  * ignored. The 'st_ino' field is ignored except if the 'use_ino'
162  * mount option is given.
163  */
164 int caching_getattr(const char *path, struct stat *statbuf){
165     writeToLog("getattr");
166
167     if(isLogFile(path)) {
168         return -ENOENT;
169     }
170
171     char fpath[PATH_MAX];
172     buildPath(fpath, path);
173     int ret = stat(fpath, statbuf);
174     if(ret != 0) {
175         return -errno;
176     }
177     return ret;
178 }
179
180 /**
181  * Get attributes from an open file
182  *
183  * This method is called instead of the getattr() method if the
184  * file information is available.
185  *
186  * Currently this is only called after the create() method if that
187  * is implemented (see above). Later it may be called for
188  * invocations of fstat() too.
189  *
190  * Introduced in version 2.5
191  */
192 int caching_fgetattr(const char *path, struct stat *statbuf,
193                     struct fuse_file_info *fi){
194     writeToLog("fgetattr");
195

```

```

196     if(isLogFile(path)) {
197         return -ENOENT;
198     }
199
200     int ret = fstat((int) fi->fh, statbuf);
201     if(ret != 0) {
202         return -errno;
203     }
204     return ret;
205 }
206
207 /**
208  * Check file access permissions
209  *
210  * This will be called for the access() system call. If the
211  * 'default_permissions' mount option is given, this method is not
212  * called.
213  *
214  * This method is not called under Linux kernel versions 2.4.x
215  *
216  * Introduced in version 2.5
217  */
218 int caching_access(const char *path, int mask)
219 {
220     writeToLog("access");
221
222     if(isLogFile(path)) {
223         return -ENOENT;
224     }
225
226     char fpath[PATH_MAX];
227     buildPath(fpath, path);
228
229     int ret = access(fpath, mask);
230     if(ret != 0) {
231         return -errno;
232     }
233     return ret;
234 }
235
236
237 /** File open operation
238  *
239  * No creation, or truncation flags (O_CREAT, O_EXCL, O_TRUNC)
240  * will be passed to open(). Open should check if the operation
241  * is permitted for the given flags. Optionally open may also
242  * initialize an arbitrary filehandle (fh) in the fuse_file_info
243  * structure, which will be passed to all file operations.
244  *
245  * pay attention that the max allowed path is PATH_MAX (in limits.h).
246  * if the path is longer, return error.
247  *
248  * Changed in version 2.2
249  */
250 int caching_open(const char *path, struct fuse_file_info *fi){
251     writeToLog("open");
252
253     if(isLogFile(path)) {
254         return -ENOENT;
255     }
256
257     fi->direct_io = 1;
258     char fpath[PATH_MAX];
259     buildPath(fpath, path);
260     // if path too long
261     if(sizeof(*fpath) > PATH_MAX) {
262         return -ENAMETOOLONG;
263     }

```



```

264
265 // check the access is valid (read only)
266 if((fi->flags & 3) != O_RDONLY) {
267     return -EACCES;
268 }
269
270 int ret = open(fpath, (O_RDONLY | O_DIRECT | O_SYNC));
271 if(ret == -1) {
272     return -errno;
273 }
274 fi->fh = (uint64_t) ret;
275 return 0;
276 }
277
278
279 /** Read data from an open file
280  *
281  * Read should return exactly the number of bytes requested except
282  * on EOF or error. For example, if you receive size=100, offset=0,
283  * but the size of the file is 10, you will init only the first
284  * ten bytes in the buff and return the number 10.
285
286  In order to read a file from the disk,
287  we strongly advise you to use "pread" rather than "read".
288  Pay attention, in pread the offset is valid as long it is
289  a multiplication of the block size.
290  More specifically, pread returns 0 for negative offset
291  and an offset after the end of the file
292  (as long as the the rest of the requirements are fulfilled).
293  You are suppose to preserve this behavior also in your implementation.
294
295  * Changed in version 2.2
296  */
297 int caching_read(const char *path, char *buf, size_t size,
298                 off_t offset, struct fuse_file_info *fi){
299     writeToLog("read");
300
301     if(isLogFile(path)) {
302         return -ENOENT;
303     }
304
305     int currentBlock = (int) offset / blockSize;
306     CDE * cde;
307     ssize_t b = 0;
308     string fileName = string(path);
309     int readTotal = 0;
310     off_t newOffset;
311     bool firstRead = true;
312     //TODO what if offset + size > file_size?
313     while(true) { // will end when pread returns 0
314         newOffset = offset + (off_t) readTotal;
315         if(newOffset%blockSize != 0 && !firstRead){
316             return readTotal;
317         }
318
319         cacheMap.empty();
320
321         if(cacheMap.count({fileName, currentBlock}) > 0) {
322             // cache hit
323             cde = cacheMap[{fileName, currentBlock}];
324
325             int readSize = 0;
326             int x = blockSize * currentBlock;
327             if(readTotal == (int) size) {
328                 return readTotal;
329             } else if((size_t) newOffset >= cde->getSize() + x) {
330                 return readTotal;
331             } else if(newOffset + size < cde->getSize() + x) {

```

```

332         readSize = (int) size;
333     } else {
334         readSize = (cde->getSize() + x) - (int) newOffset;
335     }
336
337     if((int) size <= readTotal + readSize) {
338         readSize = size - readTotal;
339     }
340
341     int inBlockOffset = (int) newOffset - currentBlock * blockSize;
342
343     memcpy(buf + readTotal, cde->getData() + inBlockOffset,
344           (size_t) readSize);
345     countChain.increment(cde);
346     lru.reinsert(cde);
347     readTotal += readSize;
348
349 } else {
350     // cache miss
351     char *blockData = (char *) aligned_alloc(blockSize,
352                                           blockSize * sizeof(char));
353     off_t tmpOffset = newOffset - newOffset%blockSize;
354     b = pread((int) fi->fh, (void *) blockData, (size_t) blockSize,
355              tmpOffset);
356
357     if (b < 0) {
358         cout << -errno << endl;
359     } else if (b == 0) {
360         free(blockData);
361         return readTotal;
362     }
363
364
365     cacheMap[{fileName, currentBlock}] = new CDE(currentBlock,
366                                                  fileName, b,
367                                                  blockData);
368     CDE *cde = cacheMap[{fileName, currentBlock}];
369
370     free(blockData);
371     // add the new cde (which has count of 1 to CountChain[0]
372     countChain.insert(cde, 1);
373     if ((int) lru.getSize() < numberOfBlocks) {
374         // there is empty place in cache
375         lru.insert(cde);
376         cde->setIsNew(true);
377     } else {
378         // cache is full, use replacement policy
379         if (numberOfBlocks != 0) {
380             removeFromCache();
381             lru.insert(cde);
382         }
383     }
384
385     int readSize = 0;
386     int x = blockSize * currentBlock;
387     if(readTotal == (int) size) {
388         return readTotal;
389     } else if((size_t) newOffset >= cde->getSize() + x) {
390         return readTotal;
391     } else if(newOffset + size < cde->getSize() + x) {
392         readSize = (int) size;
393     } else {
394         readSize = (cde->getSize() + x) - (int) newOffset;
395     }
396
397     if((int) size <= readTotal + readSize) {
398         readSize = size - readTotal;
399     }

```

```

400
401         int inBlockOffset = (int) newOffset - currentBlock * blockSize;
402
403         memcpy(buf + readTotal, cde->getData() + inBlockOffset,
404                (size_t) readSize);
405
406         readTotal += readSize;
407     }
408     currentBlock++;
409     firstRead = false;
410 }
411 }
412
413 /** Possibly flush cached data
414 *
415 * BIG NOTE: This is not equivalent to fsync(). It's not a
416 * request to sync dirty data.
417 *
418 * Flush is called on each close() of a file descriptor. So if a
419 * filesystem wants to return write errors in close() and the file
420 * has cached dirty data, this is a good place to write back data
421 * and return any errors. Since many applications ignore close()
422 * errors this is not always useful.
423 *
424 * NOTE: The flush() method may be called more than once for each
425 * open(). This happens if more than one file descriptor refers
426 * to an opened file due to dup(), dup2() or fork() calls. It is
427 * not possible to determine if a flush is final, so each flush
428 * should be treated equally. Multiple write-flush sequences are
429 * relatively rare, so this shouldn't be a problem.
430 *
431 * Filesystems shouldn't assume that flush will always be called
432 * after some writes, or that it will be called at all.
433 *
434 * Changed in version 2.2
435 */
436 int caching_flush(const char *path, struct fuse_file_info *fi)
437 {
438     writeToLog("flush");
439
440     if(isLogFile(path)) {
441         return -ENOENT;
442     }
443
444     return 0;
445 }
446
447 /** Release an open file
448 *
449 * Release is called when there are no more references to an open
450 * file: all file descriptors are closed and all memory mappings
451 * are unmapped.
452 *
453 * For every open() call there will be exactly one release() call
454 * with the same flags and file descriptor. It is possible to
455 * have a file opened more than once, in which case only the last
456 * release will mean, that no more reads/writes will happen on the
457 * file. The return value of release is ignored.
458 *
459 * Changed in version 2.2
460 */
461 int caching_release(const char *path, struct fuse_file_info *fi){
462     writeToLog("release");
463
464     if(isLogFile(path)) {
465         return -ENOENT;
466     }
467

```

```

468     int ret = close(fi->fh);
469     if(ret != 0) {
470         return -errno;
471     }
472     return ret;
473 }
474
475 /** Open directory
476  *
477  * This method should check if the open operation is permitted for
478  * this directory
479  *
480  * Introduced in version 2.3
481  */
482 int caching_opendir(const char *path, struct fuse_file_info *fi){
483     writeToLog("opendir");
484
485     if(isLogFile(path)) {
486         return -ENOENT;
487     }
488
489     DIR *dp;
490     int ret = 0;
491
492     char fpath[PATH_MAX];
493     buildPath(fpath, path);
494
495     dp = opendir(fpath);
496     if(dp == NULL) {
497         ret = -errno;
498     } else {
499         fi->fh = (intptr_t) dp;
500     }
501     return ret;
502 }
503
504 /** Read directory
505  *
506  * This supersedes the old getdir() interface. New applications
507  * should use this.
508  *
509  * The readdir implementation ignores the offset parameter, and
510  * passes zero to the filler function's offset. The filler
511  * function will not return '1' (unless an error happens), so the
512  * whole directory is read in a single readdir operation. This
513  * works just like the old getdir() method.
514  *
515  * Introduced in version 2.3
516  */
517 int caching_readdir(const char *path, void *buf,
518                    fuse_fill_dir_t filler,
519                    off_t offset, struct fuse_file_info *fi){
520     writeToLog("readdir");
521
522     if(isLogFile(path)) {
523         return -ENOENT;
524     }
525
526     int ret = 0;
527     DIR *dp;
528     struct dirent *de;
529
530     dp = (DIR *) fi->fh;
531     de = readdir(dp);
532
533     if(de == 0) {
534         return -errno;
535     }

```

```

536
537     do {
538         if(filler(buf, de->d_name, NULL, 0) != 0) {
539             return -ENOMEM;
540         }
541     } while((de = readdir(dp)) != NULL);
542
543     return ret;
544 }
545
546 /** Release directory
547  *
548  * Introduced in version 2.3
549  */
550 int caching_releasedir(const char *path, struct fuse_file_info *fi){
551     writeToLog("releasedir");
552
553     if(isLogFile(path)) {
554         return -ENOENT;
555     }
556
557     int ret = closedir((DIR *) fi->fh);
558     if(ret != 0) {
559         return -errno;
560     }
561     return ret;
562 }
563
564 /** Rename a file */
565 int caching_rename(const char *path, const char *newpath){
566     writeToLog("rename");
567
568     if(isLogFile(path)) {
569         return -ENOENT;
570     }
571
572     char fpath[PATH_MAX];
573     char fnewpath[PATH_MAX];
574
575     buildPath(fpath, path);
576     buildPath(fnewpath, newpath);
577
578     int ret = rename(fpath, fnewpath);
579     if(ret != 0) {
580         return -errno;
581     }
582
583     // rename block in cache
584     CDE * cde = lru.getHead();
585     while(cde != nullptr) {
586         string fileName = cde->getFileName();
587         int pos = fileName.find(string(path));
588         if(pos == 0 && (fileName[string(path).length()] == '/' ||
589             fileName.length() == string(path).length())) {
590
591             string suffix = fileName.substr(string(path).length(),
592                 fileName.length());
593             string realNewPath = string(newpath) + suffix;
594             cde->setFileName(realNewPath);
595             cacheMap[{realNewPath, cde->getBlockId()}] = cde;
596             cacheMap.erase({string(path), cde->getBlockId()});
597         }
598         cde = cde->getNext();
599     }
600     return ret;
601 }
602
603 /**

```

```

604  * Initialize filesystem
605  *
606  * The return value will be passed in the private_data field of
607  * fuse_context to all file operations and as a parameter to the
608  * destroy() method.
609  *
610
611  If a failure occurs in this function, do nothing (absorb the failure
612  and don't report it).
613  For your task, the function needs to return NULL always
614  (if you do something else, be sure to use the fuse_context correctly).
615  * Introduced in version 2.3
616  * Changed in version 2.6
617  */
618  void *caching_init(struct fuse_conn_info *conn){
619      writeToLog("init");
620      return NULL;
621  }
622
623
624  /**
625   * Clean up filesystem
626   *
627   * Called on filesystem exit.
628
629   If a failure occurs in this function, do nothing
630   (absorb the failure and don't report it).
631
632   * Introduced in version 2.3
633   */
634  void caching_destroy(void *userdata){
635      writeToLog("destroy");
636      logFile.close();
637  }
638
639
640  /**
641   * Ioctl from the FUSE sepc:
642   * flags will have FUSE_IOCTL_COMPAT set for 32bit ioctls in
643   * 64bit environment. The size and direction of data is
644   * determined by _IOC_*() decoding of cmd. For _IOC_NONE,
645   * data will be NULL, for _IOC_WRITE data is out area, for
646   * _IOC_READ in area and if both are set in/out area. In all
647   * non-NULL cases, the area is of _IOC_SIZE(cmd) bytes.
648   *
649   * However, in our case, this function only needs to print
650   * cache table to the log file .
651   *
652   * Introduced in version 2.8
653   */
654  int caching_ioctl (const char *, int cmd, void *arg,
655                     struct fuse_file_info *, unsigned int flags, void *data){
656      CDE * cde = lru.getTail();
657      while(cde != nullptr) {
658          writeToLog(cde->getFileName() + " " +
659                    to_string((cde->getBlockId() + 1)) + " " +
660                    to_string(cde->getCount()));
661          cde = cde->getPrev();
662      }
663      return 0;
664  }
665
666
667  // Initialise the operations.
668  // You are not supposed to change this function.
669  void init_caching_oper()
670  {
671

```

```

672     caching_oper.getattr = caching_getattr;
673     caching_oper.access = caching_access;
674     caching_oper.open = caching_open;
675     caching_oper.read = caching_read;
676     caching_oper.flush = caching_flush;
677     caching_oper.release = caching_release;
678     caching_oper.opendir = caching_opendir;
679     caching_oper.readdir = caching_readdir;
680     caching_oper.releasedir = caching_releasedir;
681     caching_oper.rename = caching_rename;
682     caching_oper.init = caching_init;
683     caching_oper.destroy = caching_destroy;
684     caching_oper.ioctl = caching_ioctl;
685     caching_oper.fgetattr = caching_fgetattr;
686
687
688     caching_oper.readlink = NULL;
689     caching_oper.getdir = NULL;
690     caching_oper.mknod = NULL;
691     caching_oper.mkdir = NULL;
692     caching_oper.unlink = NULL;
693     caching_oper.rmdir = NULL;
694     caching_oper.symlink = NULL;
695     caching_oper.link = NULL;
696     caching_oper.chmod = NULL;
697     caching_oper.chown = NULL;
698     caching_oper.truncate = NULL;
699     caching_oper.utime = NULL;
700     caching_oper.write = NULL;
701     caching_oper.statfs = NULL;
702     caching_oper.fsync = NULL;
703     caching_oper.setxattr = NULL;
704     caching_oper.getxattr = NULL;
705     caching_oper.listxattr = NULL;
706     caching_oper.removexattr = NULL;
707     caching_oper.fsyncdir = NULL;
708     caching_oper.create = NULL;
709     caching_oper.ftruncate = NULL;
710 }
711
712
713 int main(int argc, char* argv[]){
714
715     // Check input parameters
716     if(!isInputParamsValid(argc, argv)) {
717         exit(0);
718     }
719
720     init_caching_oper();
721
722     blockSize = getBlockSize();
723     rootDir = realpath(argv[1], NULL);
724
725     strcat(logPath, rootDir);
726     strcat(logPath, FILE_LOCATION);
727
728     numberOfBlocks = atoi(argv[3]);
729
730     lru.setNewIndex((int) (atof(argv[5]) * numberOfBlocks));
731     lru.setOldIndex(numberOfBlocks - (int) (atof(argv[4]) *
732         numberOfBlocks) + 1);
733
734     logFile.open(logPath, std::ios_base::app); // create/open log file
735
736     // arrange args for fuse_main call
737     argv[1] = argv[2];
738     for (int i = 2; i < (argc - 1); i++){
739         argv[i] = NULL;

```

```
740     }
741     argv[2] = (char*) "-s";
742     //argv[3] = (char*) "-f";
743     argc = 3;
744
745     int fuse_stat = fuse_main(argc, argv, &aching_oper, NULL);
746     free(rootDir);
747     return fuse_stat;
748 }
```


5 CountChain.h

```
1  //
2  // Created by root on 5/23/16.
3  //
4
5  #ifndef OS_EX4_COUNTCHAIN_H
6  #define OS_EX4_COUNTCHAIN_H
7  #include <vector>
8  #include "CDE.h"
9  #include <iostream>
10
11 using namespace std;
12
13 class CountChain {
14 private:
15     vector<pair<CDE *, CDE *>> _countChain;
16 public:
17     CountChain(int CMax);
18     void insert(CDE * cde, int pos);
19     void increment(CDE * cde);
20     void remove(CDE * cde);
21     CDE * getItemToRemove();
22     void printCountChain();
23 };
24
25
26 #endif //OS_EX4_COUNTCHAIN_H
```

6 CountChain.cpp

```
1  //
2  // Created by root on 5/23/16.
3  //
4
5  #include "CountChain.h"
6
7  CountChain::CountChain(int CMax) {
8      for(int i = 0; i < CMax; ++i) {
9          auto myPair = new pair<CDE *, CDE*>;
10         myPair->first = nullptr;
11         myPair->second = nullptr;
12         _countChain.push_back(myPair);
13     }
14 }
15
16 void CountChain::insert(CDE * cde, int pos) {
17     --pos;
18     CDE * head = _countChain.at(pos)->first;
19     CDE * tail = _countChain.at(pos)->second;
20
21     if(head != nullptr) {
22         cde->setCountNext(head);
23         head->setCountPrev(cde);
24     }
25     _countChain.at(pos)->first = cde;
26
27     if(tail == nullptr) {
28         _countChain.at(pos)->second = cde;
29     }
30 }
31
32 void CountChain::increment(CDE * cde) {
33     if(!cde->getIsNew()) { // increase only if not new
34         int count = cde->getCount();
35         if((size_t) count <= _countChain.size()) {
36             remove(cde);
37             if((size_t) count < _countChain.size()) {
38                 insert(cde, count + 1);
39             }
40         }
41         cde->setCountPrev(nullptr);
42         cde->increaseCount();
43     }
44 }
45
46 void CountChain::remove(CDE * cde) {
47     if((size_t) cde->getCount() > _countChain.size()) {
48         return;
49     }
50     if(cde->getCountPrev() != nullptr) {
51         cde->getCountPrev()->setCountNext(cde->getCountNext());
52     } else {
53         _countChain.at(cde->getCount() - 1)->first = cde->getCountNext();
54     }
55
56     if(cde->getCountNext() != nullptr) {
57         cde->getCountNext()->setCountPrev(cde->getCountPrev());
58         cde->setCountNext(nullptr);
59     } else {
```

```

60         _countChain.at(cde->getCount() - 1)->second = cde->getCountPrev();
61     }
62 }
63
64 CDE * CountChain::getItemToRemove() {
65     for(auto it : _countChain) {
66         if(it->second != nullptr) {
67             if(it->second->getIsOld()) {
68                 return it->second;
69             }
70         }
71     }
72     return nullptr;
73 }
74 }

```

7 LRUStack.h

```
1  //
2  // Created by root on 5/21/16.
3  //
4
5  #ifndef OS_EX4_LRUSTACK_H
6  #define OS_EX4_LRUSTACK_H
7
8  #include "CDE.h"
9  #include <iostream>
10
11 class LRUStack {
12 private:
13     CDE * _head;
14     CDE * _tail;
15     CDE * _newBoundary;
16     CDE * _oldBoundary;
17     size_t _size;
18     int _newIndex, _oldIndex;
19 public:
20     LRUStack();
21     ~LRUStack();
22
23     size_t getSize() { return _size; };
24     void insert(CDE * cde);
25     void remove(CDE * cde);
26     void reinsert(CDE * cde);
27
28     void setNewIndex(int newIndex) { _newIndex = newIndex; };
29     void setOldIndex(int oldIndex) { _oldIndex = oldIndex; };
30
31     void setNewBoundary(CDE * cde) { _newBoundary = cde; };
32     void setOldBoundary(CDE * cde) { _oldBoundary = cde; };
33
34     CDE * setNewBoundary() { return _newBoundary; };
35     CDE * setOldBoundary() { return _oldBoundary; };
36
37     CDE * getTail() { return _tail; };
38     CDE * getHead() { return _head; };
39
40     void printLru();
41 };
42
43
44 #endif //OS_EX4_LRUSTACK_H
```

8 LRUStack.cpp

```
1  //
2  // Created by root on 5/21/16.
3  //
4
5  #include "LRUStack.h"
6
7  using namespace std;
8
9  LRUStack::LRUStack() {
10     _size = 0;
11     _head = nullptr;
12     _tail = nullptr;
13     _newBoundary = nullptr;
14     _oldBoundary = nullptr;
15 }
16
17 LRUStack::~LRUStack() {
18     CDE * cde = _head;
19     while(cde != nullptr) {
20         CDE* tmpCde = cde;
21         cde = cde->getNext();
22         delete tmpCde;
23     }
24 }
25
26 void LRUStack::insert(CDE * newCde) {
27
28     if(_size == 0) {
29         _tail = newCde;
30     } else {
31         _head->setPrev(newCde);
32         newCde->setNext(_head);
33     }
34     _head = newCde;
35
36     ++_size;
37     if(_size == (size_t) _newIndex) {
38         _newBoundary = _tail;
39     } else if(_size == (size_t) _oldIndex) {
40         _oldBoundary = _tail;
41         _oldBoundary->setIsOld(true);
42         _oldBoundary->setIsNew(false);
43     }
44
45     if(_size > (size_t) _newIndex) {
46         _newBoundary->setIsNew(false);
47         _newBoundary = _newBoundary->getPrev();
48     }
49 };
50
51 void LRUStack::remove(CDE * cde) {
52     _oldBoundary = cde->getPrev(); // prev because insert is coming
53     _oldBoundary->setIsOld(true);
54     _oldBoundary->setIsNew(false);
55
56     if(_head == _tail) {
57         _head = nullptr;
58         _tail = nullptr;
59     } else if(cde == _head) {
```

```

60     _head = cde->getNext();
61     cde->getNext()->setPrev(nullptr);
62 } else if(cde == _tail) {
63     _tail = cde->getPrev();
64     cde->getPrev()->setNext(nullptr);
65 } else {
66     cde->getPrev()->setNext(cde->getNext());
67     cde->getNext()->setPrev(cde->getPrev());
68 }
69 --_size;
70 };
71
72 void LRUStack::reinsert(CDE * cde) {
73
74     // update new boundary
75     if(cde->getIsNew()) {
76         if(_newBoundary == cde) {
77             if(cde->getPrev() != nullptr) {
78                 _newBoundary = cde->getPrev();
79             }
80         }
81     } else {
82         if(_newBoundary->getPrev() == nullptr) {
83             _newBoundary->setIsNew(false);
84             _newBoundary = cde;
85         } else {
86             _newBoundary->setIsNew(false);
87             _newBoundary = _newBoundary->getPrev();
88         }
89     }
90
91     // update old boundary
92     if(cde->getIsOld()) {
93         _oldBoundary = _oldBoundary->getPrev(); // prev because insert is coming
94         _oldBoundary->setIsOld(true);
95         _oldBoundary->setIsNew(false);
96     }
97
98     cde->setIsOld(false);
99     cde->setIsNew(true);
100
101     if(_head != cde) {
102         if(_tail == cde) {
103             _tail = cde->getPrev();
104             _tail->setNext(nullptr);
105         } else {
106             cde->getPrev()->setNext(cde->getNext());
107             cde->getNext()->setPrev(cde->getPrev());
108         }
109         cde->setNext(_head);
110         _head->setPrev(cde);
111         _head = cde;
112         _head->setPrev(nullptr);
113     }
114 };

```

9 Makefile

```
1  CC = gcc
2  RANLIB = ranlib
3
4  LIBSRC = CachingFileSystem.cpp CDE.h CDE.cpp CountChain.h
5  LIBSRC2 = CountChain.cpp LRUStack.h LRUStack.cpp
6
7  CPPFILES = CDE.cpp LRUStack.cpp CountChain.cpp CachingFileSystem.cpp
8
9  PKGFLAGS = `pkg-config fuse --cflags --libs`
10
11  LIBOBJ = $(LIBSRC:.cpp=.o)
12
13  INCS = -I.
14  CFLAGS = -Wall -std=c++11 -g $(INCS)
15  LOADLIBES = -L./
16
17  TAR = tar
18  TARFLAGS = -cvf
19  TARNAME = ex4.tar
20  TARSRCs = $(LIBSRC) $(LIBSRC2) Makefile README
21
22  all: CDE.cpp LRUStack.cpp CountChain.cpp CachingFileSystem.cpp
23      g++ -Wall -std=c++11 $(CPPFILES) $(PKGFLAGS) -o CachingFileSystem
24
25  clean:
26      rm *.o *.tar
27
28  tar:
29      $(TAR) $(TARFLAGS) $(TARNAME) $(TARSRCs)
30
```